

Image classification model using a neural network on the CIFAR-10 dataset. we will use transfer learning with TensorFlow and a pre-trained neural network to classify images from the CIFAR-10 dataset into 10 distinct categories. Each section will cover steps such as data preprocessing, model architecture selection, and performance evaluation.

```
# Install necessary libraries
```

```
!pip install tensorflow tensorflow_hub tensorflow_datasets
```

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.17.0)
Requirement already satisfied: tensorflow_hub in /usr/local/lib/python3.10/dist-packages (0.16.1)
Requirement already satisfied: tensorflow_datasets in /usr/local/lib/python3.10/dist-packages (4.9.6)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.3.25)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: h5py>=3.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.12.1)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: ml-dtypes<0.5.0,>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.1)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.21.5)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow) (75.1.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.5.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.12.2)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.64.1)
Requirement already satisfied: tensorboard<2.18,>=2.17 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.17.0)
Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.37.0)
Requirement already satisfied: numpy<2.0.0,>=1.23.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.26.4)
Requirement already satisfied: tf-keras>=2.14.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow_hub) (2.17.0)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (8.1.7)
Requirement already satisfied: dm-tree in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (0.1.8)
Requirement already satisfied: immutabledict in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (4.2.0)
Requirement already satisfied: promise in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (2.3)
Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (5.9.5)
Requirement already satisfied: pyarrow in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (17.0.0)
Requirement already satisfied: simple-parsing in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (0.1.6)
Requirement already satisfied: tensorflow-metadata in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (1.16.1)
Requirement already satisfied: toml in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (0.10.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (4.66.6)
Requirement already satisfied: array-record>=0.5.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (0.5.1)
Requirement already satisfied: etils>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from etils[enp,epath,epy,etree]>=1.6.0; python_version<3.11 in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (1.11.0)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from etils[enp,epath,epy,etree]>=1.6.0; python_version<3.11 in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (2024.10.0)
Requirement already satisfied: importlib_resources in /usr/local/lib/python3.10/dist-packages (from etils[enp,epath,epy,etree]>=1.6.0; python_version<3.11 in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (6.4.0)
Requirement already satisfied: zipp in /usr/local/lib/python3.10/dist-packages (from etils[enp,epath,epy,etree]>=1.6.0; python_version<3.11 in /usr/local/lib/python3.10/dist-packages (from tensorflow_datasets) (3.19.0)
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (13.9.3)
Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.0.8)
Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.13.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2024.12.14)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.7)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (0.7.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.0.6)
Requirement already satisfied: docstring-parser<1.0,>=0.15 in /usr/local/lib/python3.10/dist-packages (from simple-parsing->tensorflow) (0.16.1)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tensorboard<2.18,>=2.17->tensorflow) (3.0.2)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow) (2.18.0)
Requirement already satisfied: mdurl>=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->tensorflow) (0.1.2)
```

Before starting, installed necessary libraries. Used TensorFlow for neural network operations, TensorFlow Hub for pre-trained model access, and TensorFlow Datasets to load CIFAR-10 directly

```
import tensorflow_datasets as tfds
```

```
dataset, info = tfds.load("cifar10", with_info=True, as_supervised=True)
train_ds, test_ds = dataset['train'], dataset['test']
```

```
train_ds = train_ds.take(1000) # Use 1,000 samples for faster training
test_ds = test_ds.take(500)
```

```
# Display dataset information
info
```

```
tfds.core.DatasetInfo(
  name='cifar10',
  full_name='cifar10/3.0.2',
  description="""
  The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training
  images and 10000 test images.
  """,
  homepage='https://www.cs.toronto.edu/~kriz/cifar.html',
  data_dir='/root/tensorflow_datasets/cifar10/3.0.2',
  file_format=tfrecord,
  download_size=162.17 MiB,
  dataset_size=132.40 MiB,
  features=FeaturesDict({
    'id': Text(shape=(), dtype=string),
    'image': Image(shape=(32, 32, 3), dtype=uint8),
    'label': ClassLabel(shape=(), dtype=int64, num_classes=10),
  }),
  supervised_keys=('image', 'label'),
  disable_shuffling=False,
  splits={
    'test': <SplitInfo num_examples=10000, num_shards=1>,
    'train': <SplitInfo num_examples=50000, num_shards=1>,
  },
  citation="""@TECHREPORT{Krizhevsky09learningmultiple,
  author = {Alex Krizhevsky},
  title = {Learning multiple layers of features from tiny images},
  institution = {},
  year = {2009}
  }""",
)
```

The CIFAR-10 dataset contains 60,000 color images of 10 different classes, with each image sized at 32x32 pixels. The dataset is split into 1000 training images and 500 test images.

```
import tensorflow as tf
IMG_SIZE = (224, 224) # Adjusted to match MobileNetV2 expected input size

def preprocess_image(image, label):
    image = tf.image.resize(image, IMG_SIZE) # Resize to MobileNetV2 input size
    image /= 255.0 # Normalize pixel values to [0, 1]
    return image, label

train_ds = train_ds.map(preprocess_image).batch(32).cache().prefetch(buffer_size=tf.data.AUTOTUNE)
test_ds = test_ds.map(preprocess_image).batch(32).cache().prefetch(buffer_size=tf.data.AUTOTUNE)
```

To work effectively with a pre-trained model, resized each image to 224x224 pixels, the input size required for MobileNetV2, which will use in this project. Additionally, normalizing the images by scaling pixel values between 0 and 1 helps the model learn more efficiently.

```
import tensorflow as tf
import tensorflow_hub as hub
from tensorflow.keras import layers, models

# Define the feature extractor from TensorFlow Hub
feature_extractor_layer = hub.KerasLayer(
    "https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4",
    trainable=False
)

# Use a Lambda layer with an explicit output_shape argument
inputs = layers.Input(shape=(224, 224, 3))
x = layers.Lambda(lambda img: feature_extractor_layer(img), output_shape=(1280,))(inputs)
outputs = layers.Dense(10, activation='softmax')(x)
model = models.Model(inputs, outputs)
```

Used MobileNetV2, a lightweight neural network trained on the large ImageNet dataset, which has learned useful image features. This model's architecture is efficient yet powerful, making it suitable for CIFAR-10 classification. We'll use the MobileNetV2 feature extractor layer, followed by a fully connected layer that outputs 10 probabilities, one for each class in CIFAR-10. We'll freeze the feature extractor layer, meaning it won't be trained further, which reduces computation and helps prevent overfitting.

```
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
```

To compile the model, used the Adam optimizer, known for its adaptive learning rate that balances convergence and learning efficiency. Also use the `sparse_categorical_crossentropy` loss function, suitable for multi-class classification with integer labels. This setup helps the model learn how to distinguish between the 10 classes effectively.

```
history = model.fit(train_ds, epochs=3, validation_data=test_ds)
```

```

Epoch 1/3
32/32 ————— 62s 2s/step - accuracy: 0.2208 - loss: 2.2818 - val_accuracy: 0.5960 - val_loss: 1.2552
Epoch 2/3
32/32 ————— 76s 2s/step - accuracy: 0.6867 - loss: 1.0057 - val_accuracy: 0.6940 - val_loss: 0.9202
Epoch 3/3
32/32 ————— 51s 2s/step - accuracy: 0.7849 - loss: 0.7111 - val_accuracy: 0.7260 - val_loss: 0.8245

```

Training the model will involve 3 epochs, or complete passes over the dataset. Although more epochs can sometimes improve performance, we'll keep it moderate to avoid overfitting and to observe initial model behavior on CIFAR-10. Validation on the test dataset during each epoch helps monitor model performance, allowing us to check if it generalizes well to unseen data.

```
loss, accuracy = model.evaluate(test_ds)
print(f"Test Accuracy: {accuracy:.2f}")
```

```

16/16 ————— 18s 1s/step - accuracy: 0.7217 - loss: 0.8534
Test Accuracy: 0.73

```

Once trained, we evaluate the model on the test dataset, where accuracy measures how well the model can classify unseen images.


```
import matplotlib.pyplot as plt

# Function to plot test images and predictions
def plot_predictions(model, dataset):
    plt.figure(figsize=(10, 10))
    class_names = info.features['label'].names # Class names in CIFAR-10
    for images, labels in dataset.take(1):
        predictions = model.predict(images)
        for i in range(9):
            ax = plt.subplot(3, 3, i + 1)
            plt.imshow(images[i].numpy())
            plt.title(f"Pred: {class_names[tf.argmax(predictions[i])]}, True: {class_names[labels[i]]}")
            plt.axis("off")

plot_predictions(model, test_ds)
```

1/1 5s 5s/step

Pred: frog, True: horse



Pred: airplane, True: airplane




Pred: frog, True: frog



Pred: truck, True: truck



Pred: dog, True: dog



Pred: automobile, True: automobile



Limitations:

While our model performs well on CIFAR-10, limitations exist. Transfer learning using MobileNetV2 may not perfectly capture CIFAR-10's class-specific details, as it was originally trained on ImageNet. Future improvements could involve fine-tuning, where some layers of MobileNetV2 are unfrozen and retrained on CIFAR-10, allowing the model to adapt better. Additionally, a larger model architecture could improve performance, but this would require more computational resources.



✓ Conclusion:

In this project, we saw how to classify images from the CIFAR-10 dataset using transfer learning with a pre-trained MobileNetV2 model. By utilizing this approach, we saved computation time and achieved reasonable accuracy on a multi-class dataset. This process showcases how neural networks can be effectively applied to computer vision tasks, even with more complex datasets like CIFAR-10.

Start coding or [generate](#) with AI.